Bavitha Pidugu

Eden prairie, MN| +1 651-399-9801| bavithareddy1009@gmail.com|
linkedin.com/in/bavitha-pidugu-311b12290

SUMMARY

Data-focused Data Engineer with 5+ years of experience building scalable cloud-based data platforms and data warehouses using Snowflake and Azure. Strong expertise in ADF, SQL Server, Blob Storage, Python automation, and ETL pipelines with experience supporting analytics platforms for fraud, waste, and abuse detection. Proven ability to collaborate cross-functionally and deliver reliable data solutions in regulated environments.

- Proven experience developing ETL/ELT pipelines using Python, SQL, and Spark to process structured and semi-structured data.
- Skilled in performance tuning and cost optimization of cloud data services. Experienced in implementing medallion/lake house architectures, data quality frameworks, and performance optimization.
- Adept at collaborating with cross-functional teams (product owners, analysts, backend engineers) to translate business requirements into technical solutions.
- Hands-on experience supporting production pipelines, troubleshooting issues, and ensuring data reliability.
- Strong communicator with the ability to work independently in fast-paced Agile environments
- Adept at implementing data quality, validation, and reconciliation processes to ensure accuracy and compliance.
- Comfortable working in Agile/Scrum environments and delivering projects within tight deadlines.
- Demonstrated experience supporting production SaaS data platforms and resolving pipeline failures.
- Strong analytical and problem-solving skills with attention to detail.
- Hands-on exposure to CI/CD pipelines and DevOps practices for automated deployments.
- Self-motivated professional who can work independently while contributing effectively to team goals.
- Experienced in preparing datasets for Power BI dashboards and executive reporting.
- Strong communication skills, able to explain technical concepts to non-technical stakeholders.

SKILLS

Programming	Python, SQL, PySpark
Databases & Storage	Azure SQL, Big Query, ADLS Gen2, GCP Cloud Storage
BI & Reporting	Tableau
DevOps & Tools:	Azure DevOps, Git / GitHub, CI/CD Pipelines
Reporting:	Power BI (dataset support)
Methodologies:	Agile / Scrum
Big Data & Streaming:	Apache Spark, Event Hub / Pub-Sub (basic)
Cloud Platforms	Microsoft Azure (Azure Data Factory, ADLS Gen2, Azure SQL, Data bricks, Synapse) Google Cloud Platform (BigQuery, Cloud Storage, Dataflow, Composer/Airflow)
Data Engineering:	ETL / ELT Pipelines, Data Transformation, Data Quality & Validation
Big Data Frameworks	Apache Spark / PySpark

WORK EXPERIENCE

Client: RVO Health, Minneapolis, MN, USA **Jan 2024 - present**

Role: Azure Data Engineer

- Designed and implemented end-to-end Azure data pipelines using Azure Data Factory, ADLS Gen2, Azure SQL, and Data bricks to support healthcare analytics and reporting.
- Built scalable ETL/ELT workflows using Python, SQL, and Spark (Py Spark) to ingest and transform structured and semi-structured data from multiple source systems.
- Designed secure ingestion pipelines using Azure Data Factory (ADF) with Self-Hosted Integration Runtime,
- Designed scalable ingestion, enrichment, and aggregation pipelines using Python,PySpark, Spark SQL, and Delta
- Lake, aligning batch and streaming workloads with Medallion Architecture standards.
- orchestrating on-prem extraction, trigger-based scheduling, dependency chaining, and coordinated Data bricks
- Developed scalable ETL/ELT pipelines using Python, SQL, PySpark, and Databricks.
- Created optimized data models to support business intelligence dashboards and operational reporting.
- Built and maintained Snowflake data warehouse for analytics and reporting workloads.
- Implemented data validation, quality checks, and monitoring frameworks to ensure accuracy and reliability of production pipelines.
- Collaborated closely with product owners, analysts, and backend engineers to translate business requirements into technical data

solutions.
- Supported real-time and batch ingestion pipelines and performed performance tuning and cost optimization across Azure resources.
- Implemented CI/CD pipelines using Azure DevOps and Git for automated deployment of data workflows.
- Orchestrated workflows using Apache Airflow for automated data processing.
- Implemented real-time and batch ingestion pipelines using Kafka and AWS SQS.
- Managed development tasks and tracked progress using Jira.
- Collaborated with cross-functional teams to deliver data solutions on schedule.
- Implemented data quality checks and optimized performance across pipelines.
- Participated in Agile ceremonies and contributed to sprint planning, backlog grooming, and release cycles.
- Provided production support, troubleshooting pipeline failures, and maintaining documentation and run books.

Environment:
Azure Data Factory, ADLS Gen2, Azure SQL, Databricks, Python, SQL, PySpark, Azure DevOps, Git, Power BI

Client: Magnifi Financial Credit Union, Sartell, MN, USA ***June 2022 – Dec 2023***
Role: Data Engineer

- Designed and developed scalable ETL/ELT data pipelines to ingest financial data from core banking systems into centralized data platforms using Python, SQL, and Azure Data Factory.
- Built and maintained Azure Data Lake (ADLS Gen2) structures to support reporting, analytics, and regulatory requirements.
- Developed optimized data models to support financial reporting, including transactions, accounts, loans, and customer datasets.
- Implemented data quality checks and validation frameworks to ensure accuracy, completeness, and consistency of financial data.
- Created performant SQL queries and stored procedures for reconciliation, operational reporting, and business analytics.
- Collaborated closely with business analysts, finance teams, and stakeholders to gather requirements and translate them into technical solutions.
- Integrated AWS services (S3, EC2, Lambda) for cloud-native data solutions.
- Created Tableau dashboards and reports to support business decision-making.
- Used Datadog and Splunk to monitor data pipelines, troubleshoot issues, and improve reliability.
- Supported Power BI dashboards by preparing clean, curated datasets for executive and operational reporting.
- Monitored production pipelines, resolved data issues, and performed performance tuning and optimization.
- Followed Agile methodologies, participating in sprint planning, daily standups, and release cycles.
- Maintained documentation and provided knowledge transfer to internal teams.

Environment:
Azure Data Factory, ADLS Gen2, Azure SQL, Python, SQL, Power BI, Azure DevOps, Git

Client: Optum, Hyderabad, India ***Jan 2021 – Feb 2022***
Role: GCP Data Engineer

- Designed and developed scalable data pipelines on Google Cloud Platform (GCP) using Cloud Dataflow, BigQuery, Cloud Storage, and Composer (Airflow) to support healthcare analytics.
- Built ETL/ELT workflows using Python and SQL to ingest, transform, and load structured and semi-structured healthcare data.
- Implemented optimized BigQuery data models to support reporting on clinical, claims, and operational datasets.
- Developed batch and near real-time ingestion pipelines from multiple source systems and APIs.
- Performed data cleansing, validation, and quality checks to ensure accuracy and compliance with healthcare standards.
- Tuned queries and pipelines for performance and cost optimization across GCP services.
- Worked closely with analysts and business stakeholders to translate requirements into technical solutions for dashboards and reports.
- Used Git and CI/CD pipelines to manage code deployments and version control.
- Supported production pipelines, troubleshooting failures and ensuring SLA adherence.
- Participated in Agile ceremonies including sprint planning, daily standups, and retrospectives.

Environment:
GCP (Big Query, Cloud Storage, Dataflow, Composer), Python, SQL, Airflow, Git, CI/CD

EDUCATION

Concordia University-St. Paul **2022**
Masters, Information Technology and science