

SIDDARTHA REDDY BANDA

San Francisco, CA (Open to Relocate) | 209-684-6734 | siddarthareddybanda11@gmail.com | LinkedIn

PROFESSIONAL SUMMARY

Data Engineer with 2+ years building scalable data pipelines and dimensional models on AWS/GCP, Snowflake, and BigQuery to power product analytics and business intelligence. Proven track record designing ETL workflows with PySpark and Airflow that maintain >99% SLA, implementing SCD Type 2 dimensional models for accurate cohort and churn analysis, and optimizing warehouse queries to reduce costs by 25% and dashboard latency by 60%. Collaborate cross-functionally with product, analytics, and engineering teams to deliver governed datasets, enforce data quality standards, and reduce ad-hoc requests by 35% through self-serve BI solutions.

SKILLS

Programming & Query: Python, SQL, PySpark, Spark SQL, Bash

Data Engineering & ETL: Airflow (DAGs), ETL/ELT patterns, Kafka, Dataflow/Beam, SSIS, Sqoop, CDC, Data validation, automated workflows

Cloud Platforms & Warehouses: AWS (S3, EC2, EMR, Lambda), GCP (GCS, BigQuery, Pub/Sub, Dataproc), Snowflake, Redshift

Data Modeling & Governance: Star/Snowflake schema, SCD Type 1/2, Dimensional modeling, Data quality frameworks, SLA management, Privacy & security models

Visualization & Analytics: Power BI, Tableau, Excel (Advanced), Pandas, NumPy, KPI definitions, Cohort/retention analysis, data storytelling, stakeholder reporting

DevOps & Collaboration: Git/GitHub, Jenkins, Docker, Terraform, JIRA, Confluence

AI & Automation: Generative AI (LLM-assisted workflows), automated validation, AI-assisted documentation

PROFESSIONAL EXPERIENCE

CVS Health

Data Engineer | Jul 2025 – Present

- Built and maintained scalable data pipelines supporting healthcare and business analytics using cloud data platforms and modern ETL/ELT frameworks.
 - Collaborated with product, analytics, and engineering teams to define data requirements, ensure data quality, and deliver trusted datasets.
 - Responded to time-sensitive analytical and reporting requests from cross-functional stakeholders, ensuring accurate insights to support operational decisions.
 - Implemented data validation, monitoring, and alerting to maintain high data reliability and SLA compliance.
 - Optimized warehouse queries and data models to improve performance and reduce operational costs.
 - Contributed to governance, security, and compliance standards for sensitive healthcare data.
-

NuByt, Inc.

Data Engineer | Jul 2023 – Jun 2025 | San Francisco, CA

- Conceptualized and owned data architecture for product analytics and revenue operations, designing scalable PySpark pipelines on EMR/Dataproc ingesting app logs, CRM, and billing data to S3/GCS, Snowflake/BigQuery; maintained >99% on-time SLA with automated retries and alerting.
- Authored 20+ Airflow DAGs (cron and event-driven) orchestrating ingestion, validation, and ELT workflows; reduced manual reruns by ~60% through on-call runbooks and operational best practices.
- Designed and built dimensional data models (star/snowflake schema) with SCD Type 2 for customer, product, and subscription entities, enabling accurate cohort, churn, and MRR tracking for product and finance teams.
- Implemented data quality frameworks with expectation checks (null/range/FK validation) and deduplication logic; reduced dashboard data defects by 40% quarter-over-quarter.
- Optimized high-cost queries via partitioning, clustering, pruning, and materialized views; improved key dashboard latency from ~25 minutes to <10 minutes and reduced warehouse spend by ~25%.
- Created config-driven ingestion templates that standardized quality gates and cut new-source onboarding time by ~50%, accelerating data availability for cross-functional teams.
- Published governed datasets to Power BI and Tableau with metric contracts and documentation; reduced ad-hoc data pulls by ~35% and empowered self-serve analytics for product managers and analysts.
- Collaborated with software engineers, product managers, and data scientists to define logging requirements, triage data infrastructure issues, and represent key insights visually in dashboards supporting 3M+ user base.

- Determined and implemented security models based on privacy requirements, ensuring PII safeguards and evolving governance processes across allocated data domains.
 - Mentored junior team members on ETL best practices, code reviews, and Airflow patterns, fostering skill development and knowledge sharing within the data engineering team.
 - Performed root cause analysis on data quality and reporting issues, identifying upstream pipeline gaps and implementing long-term fixes.
 - Built automated workflows to replace manual reporting processes, improving turnaround time and reliability for recurring stakeholder requests.
 - Leveraged AI-assisted tools (LLMs) to accelerate SQL development, debugging, documentation, and data validation for complex pipelines.
-

EDUCATION

Belhaven University, Jackson, MS

Master of Science, Information Technology Management | 2025