

YESASWINI KALAHASTI

Senior Data Engineer

 +1 605-585-7159 |  Yesaswinikalahasti9@gmail.com |  United States

PROFESSIONAL SUMMARY

Results-driven **Senior Data Engineer with 5+ years of experience** building scalable, secure, and high-performance data platforms across **Azure, AWS, and GCP**. Specialized in designing **end-to-end ETL/ELT pipelines, real-time streaming architectures, cloud data warehouses, and lakehouse solutions** using Azure Databricks, Snowflake, Spark, and Azure Data Factory. Strong background in **big data ecosystems, data governance, performance optimization, and DevOps automation**. Proven success delivering enterprise data solutions in **retail, finance, healthcare, and enterprise analytics** environments.

CORE SKILLS

Data Engineering: ETL/ELT Pipelines, Data Modeling, Data Lakes, Data Warehousing, Data Migration

Big Data: Spark, PySpark, Hive, Hadoop, Kafka, HDFS

Cloud Platforms: Azure, AWS, GCP

Streaming: Kafka, Azure Event Hubs, Azure Stream Analytics, AWS Kinesis

Warehousing: Snowflake, Azure Synapse, Redshift, SQL Server

Governance & Security: RBAC, Azure Purview, Encryption, Key Vault, GDPR, HIPAA

DevOps & Automation: Azure DevOps, Jenkins, Git, Docker, Kubernetes, Terraform, ARM

BI & Analytics: Power BI, Tableau, SSRS, QlikSense

TECHNICAL SKILLS

Languages: Python, SQL, Scala, Java, C#, Shell Scripting

Big Data Tools: Apache Spark, PySpark, Hive, Pig, Sqoop, Kafka, Flume, HBase

Azure: Data Factory, Databricks, ADLS Gen2, Synapse, Azure SQL, Stream Analytics, Event Hubs, Logic Apps, Purview

AWS: S3, Lambda, Kinesis, DynamoDB, API Gateway, CloudFormation

GCP: BigQuery, Dataflow, Pub/Sub, Cloud Storage

Databases: SQL Server, Oracle, MySQL, Snowflake, Cosmos DB, Cassandra, MongoDB

ETL Tools: ADF, SSIS, Talend, Informatica, dbt, NiFi

DevOps: Azure DevOps, Jenkins, GitHub, Docker, Kubernetes

IaC: Terraform, ARM Templates, Bicep

Monitoring: Azure Monitor, Application Insights, SolarWinds, Redgate

PROFESSIONAL EXPERIENCE

Senior Azure Data Engineer

HEB — Irving, TX

Jun 2025 – Present

Lead engineer in building modern enterprise data platforms supporting retail analytics, financial reporting, and operational intelligence.

Key Contributions:

- Architected and implemented **enterprise lakehouse architecture** using **Azure Data Lake Gen2, Databricks Delta Lake, and Synapse**, improving data processing performance by **40%**
- Designed **end-to-end ETL/ELT pipelines** in **Azure Data Factory and Databricks**, integrating data from SQL Server, APIs, flat files, and cloud sources
- Developed scalable transformation frameworks using **PySpark and dbt**, enabling modular and reusable data workflows
- Built **real-time streaming pipelines** using **Kafka, Event Hubs, and Azure Stream Analytics** for near real-time business insights
- Engineered **Snowflake ingestion pipelines** using **Snowpipe and COPY commands**, reducing batch load times by **35%**
- Automated infrastructure provisioning using **Terraform, ARM templates, and Azure DevOps CI/CD**
- Implemented **data governance and lineage tracking** using **Azure Purview**, ensuring regulatory compliance
- Developed Power BI dashboards integrating **Synapse, Snowflake, and Databricks** data models
- Optimized performance through **partitioning, indexing, caching, and storage tiering**, reducing cloud costs by **30%**
- Led migration of legacy SQL Server workloads to **Azure Data Lake and Synapse**

Environment: Azure Data Factory, Databricks, ADLS Gen2, Synapse, Snowflake, Kafka, PySpark, dbt, Power BI, Azure DevOps

Big Data Engineer

Accenture — Hyderabad, India
May 2022 – Aug 2023

Worked on enterprise big data modernization initiatives for global clients.

Key Contributions:

- Developed distributed data pipelines using **Spark (Scala & PySpark), Hive, and Hadoop**
- Built **real-time ingestion workflows** using **Kafka and NiFi**
- Designed and optimized **Hive data models** for high-volume analytical queries
- Created Spark jobs for large-scale data transformations improving job efficiency by **25%**
- Implemented **Sqoop pipelines** for RDBMS to Hadoop data migration
- Automated ETL job scheduling using **Shell scripting and Oozie**
- Configured and tuned Hadoop clusters for performance and stability
- Collaborated with cross-functional Agile teams to deliver scalable analytics solutions

Environment: Spark, Hive, Hadoop, Kafka, Sqoop, NiFi, Scala, Unix, Oozie

Cloud Data Warehouse Engineer

Maiden Drop — Hyderabad, India
Apr 2019 – Apr 2022

Focused on building scalable data warehouse and big data ingestion solutions.

Key Contributions:

- Designed and built **ETL pipelines** using **Spark, Hive, Sqoop, and Kafka**
- Processed **~450GB of data daily**, optimizing Spark jobs using caching and partitioning
- Developed **Spark Streaming applications** for real-time ingestion from Kafka into HDFS and NoSQL
- Built dimensional data models to support reporting and analytics use cases
- Automated workflows using **Oozie and Control-M**, increasing reliability of data pipelines
- Integrated Hadoop data with **Tableau and SSIS** for business reporting
- Created reusable ingestion frameworks for structured and semi-structured data

Environment: Hadoop, Spark, Hive, Kafka, Sqoop, Oozie, Python, Scala, Tableau

PROJECT HIGHLIGHTS

Enterprise Data Lake Migration

Migrated on-prem SQL Server and Oracle systems to Azure Data Lake using ADF and Databricks, enabling centralized analytics.

Real-Time Retail Analytics Platform

Built Kafka + Spark Streaming solution for processing live sales and inventory data streams.

Snowflake Financial Data Warehouse

Designed dimensional models and automated ingestion pipelines supporting executive financial dashboards.

Cloud Cost Optimization Program

Implemented auto-scaling and tiered storage strategies, reducing infrastructure costs by **30% annually**.

EDUCATION

Master of Science in Computer & Information Sciences

University of South Dakota, USA

Bachelor of Technology in Computer Science Engineering

Andhra University, India