

SNIGDHA PANAKANTI

snigdha@myjobemails.com | +1 (314)-220-4344 | Seattle, USA | [LinkedIn](#)

SUMMARY

Data Scientist with 3+ years of experience designing and deploying machine learning, deep learning, and NLP solutions across enterprise and cloud-native environments. Skilled in Python, R, SQL, and advanced frameworks (TensorFlow, PyTorch, Hugging Face), with proven expertise in predictive modeling, time-series forecasting, anomaly detection, and real-time analytics. Adept at building scalable data pipelines, optimizing ML workflows, and delivering actionable insights through Power BI, Tableau, and Looker. Recognized for driving measurable business outcomes, including cost savings of \$6.5M, forecast accuracy improvements of 22%, and efficiency gains exceeding 35%.

TECHNICAL SKILLS

- **Programming & Scripting:** Python, R, SQL, NoSQL, Java, Scala, C++, Bash, HTML, CSS
- **Machine Learning & Deep Learning:** Supervised & Unsupervised Learning, Random Forest, SVM, Decision Trees, Linear/Logistic Regression, XGBoost, LightGBM, CatBoost, ARIMA, LSTM, CNN, RNN, GAN, FFNN, Transfer Learning, PyTorch, TensorFlow, Scikit-learn, Keras, A/B Testing, Hyperparameter Tuning (Optuna, GridSearchCV)
- **Natural Language Processing (NLP):** NLTK, spaCy, TextBlob, Gensim, BERT, GPT, HuggingFace Transformers, Topic Modeling, Sentiment Analysis, Text Classification
- **Data Engineering, Analysis & Visualization:** Apache Spark, PySpark, Hadoop, Hive, Kafka, Apache Airflow, Snowflake, dbt, ETL Pipelines, HDFS, Pandas, NumPy, SciPy, Statsmodels, MS Excel, Matplotlib, Seaborn, Plotly, Tableau, Power BI, Looker
- **Cloud & DevOps:** AWS (S3, EC2, RDS, SageMaker, Lambda, Glue), Azure (Data Factory, Azure ML, Synapse), GCP (BigQuery, Vertex AI), Docker, Kubernetes, MLflow, Jenkins, Git, GitHub Actions, CI/CD Pipelines
- **Databases & Tools:** MySQL, PostgreSQL, MS SQL Server, MongoDB, Cassandra, Redshift, Snowflake, Parquet, Avro, Jupyter Notebook, Google Colab, VS Code, PyCharm, Eclipse, Flask, FastAPI, FlaskAPI

PROFESSIONAL EXPERIENCE

Data Scientist, Walmart

Jan 2025 – Present | Remote, USA

- Developed predictive ML models in Python and Azure ML for revenue forecasting, churn prediction, and targeted campaigns, improving forecast accuracy by 22% and boosting customer retention by 14%.
- Engineered scalable real-time ETL pipelines with Kafka, Airflow, and Azure Stream Analytics to process 5M+ daily sales events, reducing analytics latency from 3 hours to 45 seconds.
- Optimized cloud-based analytics with Snowflake and Synapse, cutting query times by 35%, supporting 200+ dashboards, and enabling executives to access secure high-performance ML-driven insights on demand.
- Designed anomaly detection and outlier monitoring system for sales patterns, identifying \$12M potential losses annually, and reducing replenishment delays for top 50 high-demand SKUs by 28%.
- Built advanced Power BI and Looker dashboards integrated with dbt models and ML predictions, increasing decision-making speed for executives by 40% and enabling real-time trend analysis.
- Led quarterly AI/ML training sessions for 120+ business users, increasing adoption of self-service predictive analytics tools by 55% and fostering data-driven culture across merchandising, finance, and supply chain.
- Partnered with cross-functional merchandising and supply chain teams to automate sales trend forecasting, reducing manual monitoring time by 20+ hours weekly and accelerating data-driven replenishment strategies.

Data Scientist, Adobe

Jan 2024 – Dec 2024 | Remote, USA

- Trained ML models on 50M+ telemetry records using Python and Azure ML to predict adoption trends, improving product growth forecasts by 18% and accelerating feature rollout decisions.
- Built NLP pipelines with spaCy and Hugging Face Transformers for customer feedback analysis, achieving 89% classification accuracy and enabling targeted experience improvements that increased user satisfaction by 12%.
- Automated ETL workflows in Azure Data Factory, integrating 10+ disparate sources, reducing manual preparation by 40%, and enabling real-time analytics across global product performance dashboards.
- Designed financial forecasting models using time-series regression and ensemble learning, reducing variance in SaaS revenue projections by 17% and optimizing budget allocation for multi-billion-dollar software portfolios.
- Conducted cloud infrastructure cost analysis leveraging ML-based anomaly detection, identifying inefficiencies that saved \$6.5M annually while improving resource utilization across engineering workloads by 22%.
- Developed interactive Power BI dashboards powered by dbt and Snowflake to track user engagement, improving KPI visibility for 150+ senior leaders and enabling faster data-driven decisions.
- Partnered with marketing analytics teams to design predictive campaign dashboards, improving ROI by 21% through optimized ad targeting, audience segmentation, and real-time performance monitoring.

Data Scientist, Genpact	Aug 2020 – Dec 2021 Hyderabad, India
- Built Python preprocessing and feature engineering scripts on 2TB+ raw data, boosting ML training efficiency by 25%, reducing manual effort, and accelerating predictive analytics readiness across enterprise client projects. - Automated SQL-based ETL pipelines for analytics inputs, improving data quality by 30%, cutting preparation time from 6 hours to 2 hours, and enabling reliable predictive model development for business operations. - Developed Power BI dashboards visualizing ML outcomes and operational KPIs, increasing stakeholder adoption by 40%, reducing reporting delays, and supporting data-driven decision-making across diverse international business clients. - Automated weekly reporting and model evaluation metrics in Excel and Python, reducing manual processing by 35% and improving delivery speed for business performance dashboards by several hours weekly. - Collaborated with senior data scientists on anomaly detection for predictive models, boosting retraining cycle efficiency by 28% and improving stability, accuracy, and performance consistency across multiple client deployments. - Documented ETL pipelines, preprocessing logic, and feature transformations, ensuring reproducibility, enabling 100% audit traceability, and facilitating efficient cross-team handoffs for analytics and ML engineering workflows across client projects.	

PROJECTS

Bail Status and Amount Prediction
- Preprocessed felony bail data and built machine learning models to predict bail status and amount, ensuring transparent model logic and stakeholder usability. - Designed dashboards that simplified complex model outputs for legal analysts and business users. - Evaluated feature importance and model performance using precision, recall, and F1 metrics to guide model refinement and deployment.
Cost Analysis, Predictions, and Forecasting
- Conducted in-depth analysis of historical cost data for budgeting and financial forecasting in a cloud service environment. - Developed predictive models to forecast costs and provide actionable insights for resource allocation decisions. - Created presentations and data narratives to communicate findings to non-technical business stakeholders.

EDUCATION

Master of Science, University of Maryland Baltimore County Data Science	Jan 2022 – Dec 2023 Baltimore, USA
Bachelors of Technology, Sreyas Institute of Engineering and Technology Electronics and Communications Engineering	July 2017 – Aug 2021 Hyderabad, India