

SUMMARY

Data Scientist with around 4 years of experience applying machine learning, statistical analysis, and generative AI techniques to solve complex business and research problems. Experienced in building end-to-end data pipelines and developing predictive, NLP, and GenAI solutions using Python, SQL, Pandas, NumPy, scikit-learn, TensorFlow, and PyTorch. Proven ability to deploy and scale models on cloud platforms including AWS and Microsoft Azure, and to translate complex datasets into actionable insights through Power BI and Tableau. Strong collaborator with experience working across industry and academic research environments.

SKILLS

Programming Languages: Python, R, SQL
Data Visualization: Tableau, Power BI, Plotly, Seaborn, Matplotlib, Excel (Pivot Tables, VLOOKUP)
Cloud Platforms: AWS (S3, Glue, Redshift, Lambda), Azure (Azure ML, Azure Data Factory)
DevOps & Tools: Git, GitHub, JIRA, Confluence
Data Science & Machine Learning: Regression, Classification, Clustering, A/B Testing, Feature Engineering, NLP & GenAI (Hugging Face, LangChain, RAG, Embeddings, OpenAI, FAISS, Pinecone, Prompt Engineering), Scikit-learn, Pandas, NumPy, Statsmodels
Big Data & ETL: PySpark, Apache Spark, Databricks, Airflow, Kafka
Databases & Warehousing: Snowflake, PostgreSQL, MySQL, MongoDB
Project Methodologies: Agile (Scrum, Kanban), Waterfall
Soft Skills: Problem Solving, Decision Making, Stakeholder Communication, Time Management, Collaboration, Attention to Detail

PROFESSIONAL EXPERIENCE

Data Scientist / Gen AI engineer, Charter Communications

Sep 2025 – Present

- Developed a GenAI document retrieval prototype using OpenAI, Prompt Engineering, and LangChain, implementing RAG pipelines that reduced analyst search time by 32% across internal telecom knowledge bases.
- Implemented semantic search pipelines using Hugging Face embeddings with FAISS, improving document relevance precision by 27% in offline retrieval benchmark evaluations.
- Deployed scalable vector storage using Pinecone, enabling sub-second similarity queries for 100K+ unstructured enterprise documents during GenAI proof-of-concept experimentation.
- Orchestrated ingestion workflows using Azure Data Factory and experimentation pipelines in Azure ML, supporting repeatable model testing across 5 internal datasets without manual data preparation.

Data Scientist / Research Assistant Saint Louis University

Jan 2024 – May 2025

- Built statistical research models using Regression with Statsmodels (Python), improving explanatory power by 18% over baseline approaches, providing actionable insights for longitudinal academic research across multiple datasets.
- Trained supervised learning pipelines using Classification algorithms in Scikit-learn, achieving 78% accuracy across controlled academic research experiments, enhancing predictive reliability for experimental and observational study outcomes.
- Conducted structured data preparation using Pandas and NumPy, curating over 50K records and storing reproducible research datasets in AWS S3, ensuring consistent, version-controlled datasets for future analysis and modeling.
- Visualized analytical findings using Tableau and Matplotlib, producing figures adopted in 3 academic reports and conference submissions, improving clarity and impact of statistical results for stakeholders.

Data Scientist, Ganit Inc

Feb 2022 - Apr 2023

- Designed predictive datasets through Feature Engineering, improving downstream model stability and reducing feature sparsity by 22% across retail datasets, enabling more accurate customer and demand prediction models.
- Performed customer segmentation using Clustering algorithms, identifying 6 actionable customer cohorts adopted by pricing and demand planning teams to enhance personalized marketing strategies and optimize business outcomes.
- Executed controlled experiments using A/B Testing, validating pricing interventions that improved conversion metrics by 5%, supporting data-driven decision-making for marketing and sales optimization strategies effectively.
- Built executive dashboards using Power BI, monitoring 20+ KPIs sourced from Snowflake, enabling leadership to review weekly performance trends and make informed strategic business decisions efficiently.
- Processed high-volume datasets using PySpark, Apache Spark, and Databricks, transforming 1 TB of raw data via AWS Glue pipelines, streamlining ETL workflows for analytics and reporting pipelines.

Taruni, Data Analyst

Jan 2021 - Jan 2022

- Authored complex analytical queries using SQL, generating recurring performance reports for 10+ stakeholders from MySQL transactional datasets, providing timely insights to support operational and business decision-making effectively.
- Analyzed sales and inventory trends using Excel Pivot Tables and VLOOKUP, reducing manual reporting effort by 30%, enabling faster, more accurate operational planning and supply chain optimization.
- Created exploratory visualizations using Seaborn, identifying data inconsistencies that improved reporting accuracy by 15%, supporting analysts in uncovering trends and validating operational dataset quality.
- Developed interactive exploratory charts using Plotly, supporting faster insight reviews during weekly operational meetings tracked in JIRA and Confluence, helping teams quickly identify patterns and business opportunities.
- Automated scheduled reporting jobs using AWS Lambda, loading summarized outputs into AWS Redshift, cutting report refresh time from multiple hours to under 10 minutes, significantly improving operational reporting efficiency.

EDUCATION

Master of Science in Information System

Saint Louis University, Missouri, United States

Aug 2023 - May 2025

Bachelor in Civil Engineering

SASTRA University, Tamil Nadu, India

Jun 2018 - Apr 2022

CERTIFICATIONS

- Azure Fundamentals (AZ 900) : Microsoft
- The Data Science Course 2023: Complete Data Science Bootcamp: Udemy
- Sparkling Star & Learning Paladin: Cognizant