# Mohit Pullawar

**DATA ENGINEER | BUSINESS INTELLIGENCE & ANALYTICS SPECIALIST**

**Chicago, IL | mohitpullawar1102@gmail.com | LinkedIn | +1 (708) 582-0535**


## PROFESSIONAL SUMMARY

Results-driven **Data Engineer with over 4 years of experience** designing, optimizing, and scaling end-to-end data solutions across **supply-chain and healthcare domains**. Proven expertise in building robust **ETL pipelines, cloud data warehouses**, and **analytics platforms** that enable real-time decision-making and operational efficiency. Adept in **Python, SQL, Snowflake, AWS, DBT**, and **Airflow**, with hands-on experience in integrating Gen AI into data workflows. Recognized for transforming complex datasets into actionable insights and collaborating cross-functionally to deliver measurable business value.

## SKILLS

| | |
|---|---|
| **Certifications:** | AWS Cloud Practitioner |
| **Languages:** | Python, R programming, SQL, SAS, MATLAB |
| **Packages:** | Pandas, NumPy, SciPy, Scikit-Learn, TensorFlow, PySpark, Matplotlib |
| **Databases:** | MySQL, PostgreSQL, Oracle, MongoDB |
| **Exploratory Data Analysis:** | Data Warehousing, Data Wrangling, Data Modelling, Data Mining, Data Visualization |
| **Data Science & Machine Learning:** | Gen AI Integration, Supervised Learning, Unsupervised Learning |
| **Bigdata Technologies:** | Hadoop, MapReduce, HDFS, Hive, Pig |
| **Tools:** | Tableau, Power BI, Looker, SharePoint, Jira, Visio, Lucidchart, Docker, Kubernetes, Terraform |
| **Statistics:** | Anova Testing, Bayesian Inference, Hypothesis Testing, A/B Testing, Elasticsearch |
| **Clouds:** | AWS (EC2, S3, R53, IAM, Redshift), Azure (Data Lake, Data factory), GCP |
| **Others:** | ETL, DBT, Airflow, Databricks, Risk Analysis, Snowflake, Decision Analysis, Financial Analysis, Documentation, Reporting, Traceability Matrices, Ad Hoc Analysis KPI Analysis |
| **Version Control & CI/CD:** | Git, Bitbucket, Jenkins |
| **Methodologies:** | SDLC, Agile, Scrum, Kanban, Waterfall |

## EXPERIENCE

**Chewy Inc., | *Business Intelligence Engineer*** **May 2024 - Present, Florida**

- Designed and implemented robust **AWS-based data pipelines** using **Glue, Lambda, and Airflow** to automate ingestion from over **12 disparate data sources**. This initiative significantly reduced manual refresh efforts by **32%** and achieved **91%** reliability in data delivery, enabling stakeholders to access timely and accurate insights.
- Architected an optimized **Snowflake data warehouse** that consolidated supplier, order, and logistics data. By enhancing partitioning and join logic, query performance improved by approximately **22%**, accelerating dashboard responsiveness and decision-making efficiency.
- Collaborated with **supply chain leadership** to design dynamic dashboards using **Power BI and Amazon Quicksight**. These dashboards provided real-time visibility into inventory turnover and supply health, leading to a **7% reduction** in excess inventory across critical SKUs.
- Partnered closely with **business stakeholders** and **data scientists** to deploy advanced demand forecasting models. Model accuracy **increased from 81% to 90%**, enabling more accurate planning and reducing stockouts in high-demand regions.
- Led a **4-member cross-functional analytics team** responsible for modernizing KPI reporting and automating legacy ETL workflows. This modernization initiative resulted in a **25%** reduction in manual reporting time and improved consistency in business metrics delivery.
- Pioneered a GenAI proof-of-concept leveraging **AWS Bedrock** and Comprehend to automatically summarize supplier performance reports. This initiative cut manual review time by nearly **50%**, increasing analyst productivity and enabling quicker issue detection.

**VowelWeb | *Data Analyst / Data Engineer*** **July 2021 - August 2023, India**

- Led the migration of outdated **SSIS ETL pipelines** to modern solutions using **Azure Data Factory** and **Synapse Analytics**. Daily data load times were reduced from **8 hours to just 2 hours**, while compute resource consumption was optimized, leading to a **20%** cost reduction.
- Designed and developed star-schema data models and automated **15+ interactive Power BI dashboards**. These dashboards were actively used by over **200 stakeholders**, significantly improving real-time reporting capabilities and enhancing decision-making across supply planning and inventory teams.
- Engineered data validation and reconciliation processes using **Python and SQL** that reduced data discrepancies by **25%** quarter over quarter. These checks ensured data consistency across multiple systems and boosted data trust among business users.
- Streamlined and automated the **generation of key warehouse performance KPIs**, contributing to a **10%** improvement in fulfillment lead times and a **5% gain** in inventory optimization.
- Collaborated with machine learning engineers to build a **supplier segmentation model** using unsupervised learning techniques. This enabled the logistics team to personalize vendor strategies and save up to **8%** in transportation and logistics costs.

**VowelWeb | *Junior Data Analyst*** **July 2020 - June 2021, India**

- Built and maintained **scalable ETL pipelines** using **Azure Data Factory and SQL Server** to process large volumes of EHR, claims, and laboratory data **(over 8 million records daily).** The resulting data platform improved data accessibility and availability by **45%**, supporting analytics initiatives across multiple departments.
- Developed and maintained interactive **Tableau** dashboards that visualized provider performance, patient care quality, and treatment outcomes. These dashboards supported care coordination improvements, resulting in a **12%** increase in service efficiency.
- Created an automated **QA** testing framework in **Python** to validate **HIPAA-compliant reports**. This saved the team approximately **25 hours per month** in manual testing and reduced the risk of regulatory compliance issues.
- Built a machine learning model using Logistic Regression and Random Forest techniques to predict patient readmission risk. The model achieved an AUC of **0.81** and contributed to an **8%** decrease in readmission rates by helping care teams proactively intervene.

## EDUCATION

**Masters of Science in Business Analytics** **GPA: 3.93**
DePaul University, Chicago Sep 2023 – Jun 2025

**Bachelors of Engineering in Electrical Engineering** **GPA: 3.54**
Savitribai Phule Pune University, India Jun 2017 – Jun 2021

## PROJECT

**Retail Demand Forecasting Platform | Snowflake, Airflow, Python, XGBoost, Power BI**

- Built a data pipeline using Airflow for ingesting POS and inventory data into Snowflake and structured it for analytical processing.
- Cleaned and preprocessed time-series data in Python (Pandas, NumPy), preparing it for regression-based demand forecasting.
- Developed and fine-tuned an XGBoost model to forecast product-level sales, incorporating seasonality and store-level variation.
- Visualized forecast accuracy and key KPIs in Power BI to support retail demand planning.

**Healthcare Analytics Automation Suite | ADF, Power BI REST API, Python, SQL**

- Engineered a fully automated Azure Data Factory pipeline integrating EHR, claims, and patient satisfaction data.
- Leveraged Power BI REST API to schedule dashboard refreshes and ensure real-time reporting for stakeholders.
- Created a Python-SQL validation module to verify patient-provider records before reporting.
- Delivered Power BI dashboards highlighting key metrics like provider efficiency, care quality, and patient sentiment.

**Unified Customer 360 Platform | DBT, Redshift, AWS Glue, Tableau**

- Modeled multi-source customer data using DBT and centralized it in Redshift for advanced segmentation and LTV analysis.
- Automated data ingestion via AWS Glue from CRM, marketing, and support platforms into a unified data mart.
- Built Tableau dashboards showcasing behavioral insights, retention metrics, and omni-channel engagement trends.