

Rehaan Arya

✉ rehaanarya@berkeley.edu || ☎ 972-971-5865 || in /rehaan-arya

EDUCATION

University of California at Berkeley

Berkeley, California

Bachelor's in Data Science

August 2022 - December 2025

- **Concentrations:** Economics, Business and Industrial Analytics
- **Related Coursework:** Principles and Techniques of Data Science, Data Engineering, Data Mining and Analytics, Data & Decisions, Computer Programming, Data Structures, Probability for Data Science, Linear Algebra and Differential Equations, Generative AI, Technology Entrepreneurship

SKILLS

Data Analysis & Visualization: Tableau (LOD expressions, parameters, drill-down dashboards), Power BI, Looker (exposure), Advanced Excel (Pivot Tables, VLOOKUP, INDEX-MATCH), Matplotlib, Seaborn, Plotly

SQL & Databases: SQL (CTEs, window functions, large-scale enterprise databases), PostgreSQL, MySQL, MS SQL Server

Machine Learning & Modeling: Scikit-learn (Random Forest, K-Nearest Neighbors), TPOT, Pandas, NumPy, Statsmodels, TensorFlow, PyTorch

Tools & Technologies: Jupyter Notebooks, Git, VS Code, IntelliJ, AWS (Lambda, Step Functions, S3, SQS, OpenSearch, DynamoDB, Bedrock, EC2), PySpark, Snowpark

PROFESSIONAL EXPERIENCE

Amazon Web Services - Annapurna Labs

Cupertino, California

Machine Learning Engineer

September 2025 – December 2025

- Architected and deployed end-to-end AWS data ingestion and ETL pipelines (Lambda, Step Functions, S3, SQS) to process 1,200+ research papers in under two minutes for downstream machine learning evaluation
- Designed scalable model evaluation and training workflows in PyTorch, enabling rapid experimentation and benchmarking of Trainium-compatible techniques
- Collaborated with cross-functional research and infrastructure teams to surface performance metrics and operational insights from large-scale ML experiments

Capgemini

New York City, New York

Data Science and Machine Learning Intern

June 2025 – August 2025

- Performed exploratory data analysis and advanced visualizations (Seaborn, Matplotlib) to analyze user query behavior and improve LLM conversational accuracy
- Integrated model-based microservices into a real-time AI pipeline powering a digital financial advisor
- Used Python to test and fine-tune speech and language models, reducing latency and API response times by 25%

Walmart Global Tech

Bentonville, Arkansas

Tech Strategy Consultant

January 2025 – May 2025

- Designed, built, and maintained dynamic Tableau dashboards using parameters and drill-down visualizations to monitor application and workforce performance, contributing to a 10% improvement in operational efficiency
- Conducted in-depth exploratory data analysis on large enterprise datasets using SQL (CTEs, window functions) and Python (Pandas, NumPy), transforming raw data into actionable insights
- Presented analytical findings and strategic recommendations to cross-functional leadership, identifying \$500K in potential savings through automation and generative AI initiatives

Applied Materials

Santa Clara, California

Data Scientist

September 2024 – December 2024

- Built scalable data pipelines to extract and preprocess semiconductor performance data from unstructured sources using OCR technologies
- Designed and deployed predictive modeling pipelines leveraging tree-based and ensemble methods to forecast semiconductor behavior with 85% accuracy
- Automated data preprocessing, feature engineering, and model training workflows to support reproducible experimentation and data-driven decision-making

PROJECTS

Enterprise Sales & Operations Analytics Dashboard

January 2025 - May 2025

- Built automated ETL workflows in Python and SQL to ingest and transform transactional data for recurring analysis and reporting
- Designed interactive Tableau dashboards using LOD expressions, parameters, and drill-down views to track KPIs and operational performance
- Delivered recurring insights that supported data-informed decision-making and business growth initiatives

Predictive Demand Forecasting Pipeline

September 2024 - December 2024

- Developed end-to-end predictive modeling pipelines using scikit-learn, implementing Random Forest and K-Nearest Neighbors algorithms
- Automated model selection and hyperparameter tuning using TPOT to improve prediction accuracy and model efficiency
- Evaluated model performance using cross-validation and business-aligned metrics, translating results into actionable recommendations

March ML Madness

January 2024 - March 2024

- Trained and evaluated machine learning models using scikit-learn and TensorFlow to predict NCAA basketball outcomes with 80.3% accuracy
- Performed feature engineering, cross-validation, and automated hyperparameter tuning to enhance model reliability and performance