

Praneeth Dala

Praneeth9924@gmail.com | +1-659-253-9854 | Birmingham, AL | [LinkedIn](#)

Summary

Results-driven Data & Cloud Engineer with 4+ years of experience designing and scaling cloud-native data platforms across AWS and GCP. Specialized in ETL development, real-time data streaming, data warehousing, and analytics enablement. Proven success in reducing data latency by 70%, improving ML retraining frequency by 25%, and delivering 99.9% production availability. Proficient in Python, SQL, Kubernetes, Terraform, and CI/CD workflows. Open to Data Engineering, Data Analytics, Software Engineering, and Cloud Engineering opportunities.

Core Skills

Programming & Data: Python, SQL, PySpark, Apache Spark, Go

Cloud Platforms: AWS (Glue, Lambda, Redshift), GCP (BigQuery), Azure (Databricks, Data Factory)

ETL & Data Pipelines: Kafka, Airflow, Hadoop, Kinesis, Data Lake Storage

DevOps & Infrastructure: Terraform, GitLab CI/CD, Jenkins, Docker, Kubernetes (EKS/GKE, Helm)

Data Analytics & ML: Power BI, Tableau, Looker, Scikit-Learn, TensorFlow, Predictive Modeling

Database & Warehousing: Snowflake, Redshift, BigQuery, PostgreSQL, MySQL

Other: Data Modeling, Data Governance, DataOps, Cloud Security, CI/CD Automation

Professional Experience

Citi Group - New York, NY | July 2024 – May 2025

Data Engineer

- Implemented real-time fraud detection pipelines using AWS Glue, Lambda, Kinesis, and Redshift, processing 5,000+ endpoints with sub-minute latency.
- Deployed and maintained Kubernetes workloads on EKS using Helm and CRDs, achieving 99.9% production availability.
- Automated infrastructure provisioning with Terraform, reducing manual setup time by 70% and enabling reusable environments.
- Implemented CI/CD and GitOps workflows for Glue jobs and Redshift updates, minimizing release errors and accelerating delivery by 40%.
- Optimized data warehousing and analytics pipelines in Snowflake, reducing query runtime by 30%.
- Collaborated with fraud analytics teams to deliver ML-ready datasets, increasing model retraining frequency by 25%.

Tech Stack: AWS Glue, Lambda, Redshift, Kinesis, Kubernetes (EKS), Helm, Terraform, GitLab CI/CD, Python, SQL

QBE - Dallas, TX | May 2023 – May 2024

Data Engineer

- Designed serverless ETL pipelines with AWS Glue and Lambda, cutting fraud detection latency from hours to minutes for high-priority transactions.
- Developed near real-time ingestion pipelines using Kinesis and Spark Structured Streaming to process millions of daily credit card transactions.
- Built reusable PySpark transformation packages on Databricks, improving data reusability and development speed by 25%.
- Integrated CloudWatch and Datadog monitoring to proactively detect issues, reducing pipeline outages by 35%.
- Partnered with compliance teams to implement schema validation and audit logging to meet SOX and GDPR requirements.

Tech Stack: AWS Glue, Lambda, Kinesis, Spark Structured Streaming, PySpark, Databricks, CloudWatch, Datadog, Python, SQL

Wipro Technologies Limited - Bangalore, India | Aug 2020 – Nov 2022

Project Engineer

- Delivered enterprise-wide ingestion and transformation pipelines using AWS Glue/Redshift and Azure Synapse/Data Factory for financial analytics.
- Developed Kafka-based Java microservices to process 100,000+ messages per second, ensuring scalable event-driven architectures.
- Enhanced backend CLI tools in Go, automating operational workflows and cutting manual tasks by 50%.
- Standardized governance and validation using Terraform, improving deployment reliability and reducing errors by 40%.
- Mentored junior engineers in AWS and Spark best practices, increasing team autonomy and efficiency.

Tech Stack: AWS Glue, Redshift, Azure Synapse, Data Factory, Kafka, Java, Go, Terraform, Spark

Key Achievements

- Deployed real-time fraud detection pipelines reducing latency from hours to <1 minute for 5,000+ endpoints.
- Automated infrastructure provisioning with Terraform, cutting setup time by 70%.
- Enabled 99.9% uptime for production Kubernetes workloads with Helm and CRDs.
- Delivered ML-ready datasets that boosted retraining cycles by 25%.
- Optimized monitoring and alerting pipelines, reducing outage incidents by 35%.

Projects

Predictive Analytics ML Pipeline

Built an end-to-end PySpark and SageMaker ML pipeline for financial datasets, cutting experimentation time by 30% and improving decision accuracy by 25%.

Tech Stack: PySpark, AWS SageMaker, Scikit-Learn, Pandas, Python, REST APIs

Remote Data Integrity Verification

Implemented a cryptographic verification protocol using Merkle trees and key generation, reducing bandwidth overhead by 40% and ensuring secure third- party integrity checks.

Tech Stack: Python, Cryptography, Hashing, Merkle Trees

Education & Certification

University of Alabama at Birmingham — Birmingham, AL | Jan 2023 – May 2024

M.S. in Computer Science

Certification:

- AWS Certified Cloud Practitioner
- AWS Certified Solutions Architect – Associate
- AWS Certified Machine Learning – Specialty
- Google Cloud Professional Data Engineer

Publications

“Remote Data Integrity Checking with Data Privacy for Cloud Storage” — Journal of Xidian University, Vol. 14, Issue 5, 2020. [View Publication](#)

Professional Closure

Available to join immediately. Open to relocation across the USA. References available upon request.