

SIDDHARTH KUSHAWAHA

Jersey City, NJ | Open to Relocation

Phone: (551) 260-8978 | Email: kushb35@gmail.com

LinkedIn: <https://www.linkedin.com/in/siddharth-k-240a1812b/> | GitHub: <https://github.com/sidkush>

Data Scientist with **3+ years** building scalable data products and ML systems across NLP, Vision, and **LLMs**. Strong ownership of end-to-end data lifecycle: ingest, transform, validate, model, and serve. Expert in **SQL (Snowflake, PostgreSQL), Python, dbt, Airflow, and modern MLOps**. Proven track record optimizing queries, operationalizing data products with documentation/monitoring, and collaborating with cross-functional teams to drive data-driven decisions at scale.

EXPERIENCE

Community Dreams Foundation | Data Analyst | Sebring, FL

Sept 2025– Present

- Built and maintained production-grade data pipelines in Snowflake and AWS RDS using Python and Airflow, with dbt models for transformations and data contracts; reduced end-to-end pipeline latency by 60%.
- Implemented performance management system: columnar clustering, partition pruning, materialized views, and SQL rewrites; improved key dashboard query times by 60% and stabilized SLAs.
- Designed automated data quality checks (dbt tests + Great Expectations) with change alerts via Slack/Email; cut data incidents and ad-hoc reporting requests by 50% via self-service data definitions and documentation.
- Created a governed data sharing layer for stakeholders with access controls and semantic documentation; improved transparency in metric definitions and lineage using dbt docs and catalog.

Factspan Analytics | Senior Data Engineer | Bangalore, India

Jun 2022– Jun 2023

- Architected scalable LLM inference service on Kubernetes and vLLM with autoscaling and GPU optimization; instrumented benchmarking for latency, throughput, and token/sec; improved P95 latency while cutting cost.
- Built data ingestion and transformation pipelines for customer analytics in Snowflake and PostgreSQL; standardized schemas, metadata, and SCD handling; improved model training freshness and reliability.
- Operationalized data products with documentation, lineage, and monitoring; added alerting for drift/anomalies across RAG pipelines; integrated AWS Secrets Manager and Okta for secure access.
- Optimized customer-facing assistants with QLoRA on LLaMA-3.1-7B; reduced inference latency by 30% and maintained response quality.

Factspan Analytics | Data Engineer | Bangalore, India

Jan 2021 – Jun 2022

- Designed A/B tests and simulations (Monte Carlo) to evaluate sales campaign strategies; increased conversion rates by 15% and materially improved anomaly detection (74% improvement in data quality issue detection).
- Automated client segmentation pipelines with MLflow-based model registry, versioning, and rollback; ensured reproducibility and safe deployments.
- Built dashboards (Tableau/Looker) surfacing CTR, ROAS, LTV across >100 channels; enabled budget reallocations that improved ROI by 22%.

Factspan Analytics | Data Analyst | Bangalore, India

Jan 2020 – Dec 2021

- Led data migration across Snowflake servers; automated SQL workflows in Python; cut execution time by 70% and improved pipeline reliability.
- Established performance baselines and optimized transformation modules, improving throughput by 37% and lowering compute cost.
- Partnered with business and vendors to acquire and normalize external raw datasets, establishing data contracts, SLAs, and schema validation.

SKILLS

Data Engineering: SQL (Snowflake, PostgreSQL), Python, dbt, Airflow, Spark/PySpark, Data Modeling, ETL/ELT, CDC, Incremental models

Analytics & Experimentation: A/B testing, Simulation/Monte Carlo, Causal inference basics, Metric design, Statistical testing

Machine Learning: Model development, validation, and optimization (XGBoost/LightGBM/CatBoost, CV, hyperparameter tuning), LLM

Reliability & Observability: Data quality checks, SLAs, anomaly detection, alerts (Great Expectations/DBT tests/Monte Carlo-style checks)

Performance: Query tuning, indexing/partitioning, cost optimization, vLLM/TensorRT for inference performance

Platforms & Tools: AWS, GCP (BigQuery/Vertex), Azure Databricks, Snowflake, MLflow, Docker, Kubernetes, Git, Tableau/Looker/Power BI

EDUCATION

Stevens Institute of Technology | Hoboken, NJ | Masters in Data Science (Concentration in Business Analytics)

[GPA: 3.87/4]

Relevant Coursework - Gen AI, Probability, Statistics, Optimization, DL, ML, Linear Algebra, Web Mining, Big Data

ACADEMIC PROJECTS

[AI Investment Strategist](#) - Leveraged streamlit, yFinance, plotly, and Google Gemini LLMs to build an AI-powered platform that delivers real-time, user-driven stock analysis and personalized investment reports, enhancing research accuracy and user engagement.

[Dynamic Product Price Optimization using Amazon Product Reviews \(Big Data\)](#) - Developed a data-driven (50GB) model by analyzing customer reviews, ratings, and sentiment to identify patterns that correlate with price sensitivity and purchasing behavior with two stage sequential model – Sentiment Analysis and Price Prediction (Reinforcement), offering optimized pricing recommendations that adjust to market conditions and customer Sentiment.

ACHIEVEMENTS

Awarded "Earn Trust", "Ownership" and promoted for expanding project portfolio and revenue by 25%.