Sravya Koppula Data Engineer

sravya@mymailsbox.net | (551) 200-4486 | USA | LinkedIn

Summary

Experienced Data Engineer with 3 years of experience in designing, building, and optimizing scalable data pipelines and lakehouse architectures. Proficient in both batch and real-time data processing using a variety of modern tools and technologies. Strong skills in data quality running, monitoring, and observability to ensure high data reliability and accuracy. Adept at working across cloud platforms like AWS and Azure, with a solid understanding of data storage, transformation, and analytics to drive business insights and decision- making. Experienced with containerization, CI/CD automation, and cloud-native services to enable efficient, scalable data workflows.

Technical Skills

- **Data Engineering & Processing:** Apache Kafka, Apache Flink, Kafka Streams, Apache NiFi, Apache Spark (PySpark), dbt, SQL, Apache Beam, Google Dataflow
- **Programming & Scripting Languages:** Python (including Pandas, NumPy), Shell scripting, Java, Scala, C/C++ (familiarity)
- **Data Pipelines & ETL:** AWS Glue, Apache Airflow, Kinesis Data Firehose, Azure Data Factory, Talend, Informatica
- **Data Storage & Lakehouse Technologies:** Delta Lake, Apache Iceberg, Apache Hudi, Hive, AWS S3, Azure Data Lake Gen2, Amazon Redshift, Redshift Spectrum, Snowflake, Google BigQuery, Apache Parquet, ORC, NoSQL databases (MongoDB, DynamoDB - working knowledge)
- **Data Quality & Validation:** Great Expectations, Amazon Deequ, AWS Glue Data Catalog, schema validation, Monte Carlo, Datafold
- **Monitoring & Observability:** Datadog, OpenTelemetry, CloudWatch, Azure Monitor, Prometheus, Grafana, ELK Stack (Elasticsearch, Logstash, Kibana)
- **Cloud Platforms:** Amazon Web Services (AWS), Microsoft Azure, Google Cloud Platform (GCP)
- **Data Warehousing:** Amazon Redshift, Snowflake, Google BigQuery, Microsoft Azure Synapse Analytics, Teradata, Oracle Exadata
- **Containerization & Orchestration:** Docker, Kubernetes
- **CI/CD & Automation:** Jenkins, GitLab CI, Git, GitHub, Software Development Life Cycle (SDLC) including testing, documentation, delivery, and support
- **Cloud-Native Services:** AWS Lambda, AWS Kinesis, Azure Functions
- **Machine Learning Frameworks:** Familiarity with TensorFlow, Scikit-Learn
- **Other Tools & Concepts:** Data modeling, Data versioning, Real-time analytics, Schema evolution, Time travel queries, KPI definition, Metadata management, Columnar storage, Data lineage, Encryption & IAM

Professional Experience

Data Engineer, KPMG — 02/2024 – Present | Remote, USA

- Worked with stakeholders and client engineers during requirement gathering sessions to design a real-time data lakehouse using Apache Kafka, Delta Lake, and AWS Glue, enabling scalable ingestion across daily records.
- Engineered streaming data pipelines with Apache Flink, Kafka Streams, and Kinesis Data Firehose, supporting real-time ingestion from financial systems and enhancing processing throughput across high-velocity data streams.
- Developed modular transformation workflows using PySpark, dbt, and Apache Airflow, reducing batch job runtimes by 50% and ensuring consistent, testable logic for real-time and historical data alignment.
- Optimized storage and analytical performance by implementing Apache Iceberg and AWS S3, enabling time travel queries, schema evolution, and improving Redshift Spectrum query latency by over 70%.
- Implemented observability and monitoring with Datadog, OpenTelemetry, and CloudWatch, reducing incident resolution time by 65% and providing visibility into pipeline health and lag metrics.
- Enforced data quality controls using Great Expectations, Amazon Deequ, and schema validation with AWS Glue Data Catalog, achieving 99.9% accuracy and minimizing downstream analytics issues for internal reporting teams.

Data Engineer, Boston Consulting Group — 01/2022 – 08/2023 | Hyderabad, India

- Collaborated with client stakeholders to gather technical requirements and designed a real-time data pipeline using Kafka, Flink, and NiFi to process over millions daily events with low latency for analytics use cases.
- Developed scalable ETL pipelines using PySpark and SQL to transform streaming data and stored enriched outputs in Snowflake and Redshift to support reporting, dashboarding, and machine learning models across client environments.
- Managed data versioning and storage using Apache Hudi and Hive on Azure Data Lake Gen2 to ensure reliable retention of raw and processed data while meeting enterprise standards for compliance and traceability.
- Worked with cross-functional BCG teams to define key metrics and KPIs which improved data coverage by 45% and reduced dashboard refresh times by 30% for product and performance analytics.
- Enabled full pipeline observability using Azure Monitor, Prometheus, and Grafana and implemented automated data validation with Great Expectations to maintain schema consistency and ensure 99.9% accuracy across real-time insights platforms.

Education

Montclair State University – Montclair, New Jersey

Master in Business Analytics — 08/2023 – 12/2024